

The following was requested:

1. We take the same original picture for all attempts.
2. We take the same "noisy picture" for all attempts - and we give it MSE and SSIM score.
3. We denoise with original "float" - give MSE and SSIM.
4. We denoise with all types of quantization that you tried and give MSE and SSIM.
5. Please give a conclusion in the end - if some types of quantization are better than others.

```
In [1]: %matplotlib notebook
```

```
import argparse
import time
import os
```

```
In [2]: import matplotlib.pyplot as plt
from data import NoisyBSDSDataset
from argument import Args
from model import DnCNN, UDnCNN, DUDnCNN, QDUDnCNN
import nntools as nt
#from utils import DenoisingStatsManager, plot, NoisyBSDSDataset
import utils
import utils2
from skimage import metrics

import cv2
import numpy as np
import torch
```

```
In [3]: import torch.utils.data as td
import torch.quantization.quantize_fx as quantize_fx
from torch.quantization.fuse_modules import fuse_known_modules
import torch.nn as nn
import torch.nn.functional as F
import torch.utils.data as td
import torchvision as tv
from PIL import Image
import model
import inference
from prettytable import PrettyTable
```

```

In [4]: def main(img_set,model_path,qat=False,quantize='none'):
    args = Args()
    args.quantize = quantize
    device = 'cpu'
    if qat:
        denoise, qat_state = inference.load_model(model_path, QDUDnCNN,
args.D, args.C, device=device)
    else:
        denoise, qat_state = inference.load_model(model_path, DUDnCNN, a
args.D, args.C, device=device)

    img = []
    titles = ['clean', 'noise', 'denoise']
    x, clean = img_set
    x = x.unsqueeze(0).to(device)
    img.append(clean)
    img.append(x[0])

    if args.quantize:
        print('Quantize model...' + quantize)
        denoise = utils2.quantize_model(args.quantize, denoise, input_ex
ample=x, qat_state=qat_state)

    t = time.time()
    with torch.no_grad():
        y = denoise(x)
    img.append(y[0])

    print(f'Elapsed: {(time.time() - t) * 1000:.2f}ms')
    print(f'Image size is {x[0].shape}.')
    y1 = utils2.myimret(img[0])
    y2 = utils2.myimret(img[1])
    y3 = utils2.myimret(img[2])
    utils2.compare_images(y1,y3)

    fig, axes = plt.subplots(ncols=3, figsize=(9,5), sharex='all', share
y='all')
    for i in range(len(img)):
        utils2.myimshow(img[i], ax=axes[i])
        axes[i].set_title(f'{titles[i]}')

```

```
In [5]: def eval_origin(img_set, display=False):  
        device = 'cpu'  
  
        img = []  
        x, clean = img_set  
        # x = x.unsqueeze(0).to(device)  
  
        img.append(clean)  
        # img.append(x[0])  
        img.append(x)  
  
        y1 = utils2.myimret(img[0])  
        y2 = utils2.myimret(img[1])  
        mse, ssim = utils2.compare_images(y1, y1, display)  
        return mse, ssim
```

```
In [6]: def eval_noisy(img_set, display=False):  
        device = 'cpu'  
  
        img = []  
        x, clean = img_set  
        # x = x.unsqueeze(0).to(device)  
  
        img.append(clean)  
        # img.append(x[0])  
        img.append(x)  
  
        y1 = utils2.myimret(img[0])  
        y2 = utils2.myimret(img[1])  
        mse, ssim = utils2.compare_images(y1, y2, display)  
        return mse, ssim
```

```

In [7]: def eval_image_set(img_set,model_path,qat=False,quantize='none',display=
False,print_out=False):
    args = Args()
    args.quantize = quantize
    device = 'cpu'
    if qat:
        denoise, qat_state = inference.load_model(model_path, QDUDnCNN,
args.D, args.C, device=device)
    else:
        denoise, qat_state = inference.load_model(model_path, DUDnCNN, a
args.D, args.C, device=device)

    img = []
    titles = ['clean', 'noise', 'denoise']
    x, clean = img_set
    x = x.unsqueeze(0).to(device)

    img.append(clean)
    img.append(x[0])

    if args.quantize:
        denoise = utils2.quantize_model(args.quantize, denoise, input_ex
ample=x, qat_state=qat_state)

    t = time.time()
    with torch.no_grad():
        y = denoise(x)
    t_elapsed = (time.time() - t)*1000
    img.append(y[0])
    y1 = utils2.myimret(img[0])
    y2 = utils2.myimret(img[1])
    y3 = utils2.myimret(img[2])
    mse0, ssim0 = utils2.compare_images(y1,y2,print_out)
    mse, ssim = utils2.compare_images(y1,y3,print_out)

    if display:
        fig, axes = plt.subplots(ncols=3, figsize=(9,5), sharex='all', s
harey='all')
        for i in range(len(img)):
            utils2.myimshow(img[i], ax=axes[i])
            axes[i].set_title(f'{titles[i]}')
    return t_elapsed, quantize, mse, ssim

```

```

In [8]: def experiment(img,display=True):
    img_set = img

    t = PrettyTable(['training','quantization','time','mse','ssim'])
    t.float_format=".5"
    t.float_format['time'] = ".2"
    t.border=True
    t.align="r"
    t.align["quantization"]="l"
    t.align["training"]="l"

    path = os.environ.get('TRAINING_DIR')+'/no_qat/checkpoint.pth'
    mse_origin,ssim_origin = eval_origin(img_set, False)
    mse_noisy, ssim_noisy = eval_noisy(img_set, False)
    t.add_row(['NA (origin)', "NA", "NA",mse_origin,ssim_origin])
    t.add_row(['NA (noisy)', "NA", "NA",mse_noisy,ssim_noisy])

    eval_time_none,q, mse, ssim = eval_image_set(img_set, path, False,'none', display)
    t.add_row(['standard',q,(eval_time_none/eval_time_none)*100,mse,ssim])

    eval_time,q, mse, ssim = eval_image_set(img_set, path, False,'fx_static', False)
    t.add_row(['standard',q,(eval_time/eval_time_none)*100,mse,ssim])
    eval_time,q, mse, ssim = eval_image_set(img_set, path, False,'fx_dynamic', False)
    t.add_row(['standard',q,(eval_time/eval_time_none)*100,mse,ssim])

    path = os.environ.get('TRAINING_DIR')+'/qat/checkpoint.pth'
    eval_time,q, mse, ssim = eval_image_set(img_set, path, False,'none', False)
    t.add_row(['QAT',q,(eval_time/eval_time_none)*100,mse,ssim])
    eval_time,q, mse, ssim = eval_image_set(img_set, path, True,'fx_static', False)
    t.add_row(['QAT',q,(eval_time/eval_time_none)*100,mse,ssim])
    eval_time,q, mse, ssim = eval_image_set(img_set, path, True,'fx_dynamic', False)
    t.add_row(['QAT',q,(eval_time/eval_time_none)*100,mse,ssim])
    path = os.environ.get('TRAINING_DIR')+'/qat_fx_static/checkpoint.pth'
    # eval_time,q, mse, ssim = eval_image_set(img_set, path, False,'none', False)
    # t.add_row(['QAT_fc_static',q,(eval_time/eval_time_none)*100,mse,ssim])
    eval_time,q, mse, ssim = eval_image_set(img_set, path, True,'fx_static', False)
    t.add_row(['QAT_fc_static',q,(eval_time/eval_time_none)*100,mse,ssim])
    # eval_time,q, mse, ssim = eval_image_set(img_set, path, True,'fx_dynamic', False)
    # t.add_row(['QAT_fc_static',q,(eval_time/eval_time_none)*100,mse,ssim])
    print(t)

```

```

In [9]: def experiment_test(img,display=False):
        img_set = img

        eval_noisy(img_set, display)
        eval_noisy(img_set, display)
        eval_noisy(img_set, display)

        path = os.environ.get('TRAINING_DIR')+'/no_qat/checkpoint.pth'
        eval_image_set(img_set, path, False,'none', display)

    def a(img,display=False):
        path = os.environ.get('TRAINING_DIR')+'/no_qat/checkpoint.pth'
        eval_image_set(img_set, path, False,'none', display)
        eval_image_set(img_set, path, False,'fx_static', False)
        eval_image_set(img_set, path, False,'fx_dynamic', False)

        path = os.environ.get('TRAINING_DIR')+'/qat/checkpoint.pth'
        eval_time,q, mse, ssim = eval_image_set(img_set, path, False,'none',
False)
        eval_time,q, mse, ssim = eval_image_set(img_set, path, True,'fx_stat
ic', False)
        eval_time,q, mse, ssim = eval_image_set(img_set, path, True,'fx_dyna
mic', False)

        path = os.environ.get('TRAINING_DIR')+'/qat_fx_static/checkpoint.pt
h'
        eval_time,q, mse, ssim = eval_image_set(img_set, path, False,'none',
False)
        eval_time,q, mse, ssim = eval_image_set(img_set, path, True,'fx_stat
ic', False)
        eval_time,q, mse, ssim = eval_image_set(img_set, path, True,'fx_dyna
mic', False)

```

```

In [10]: def experiment_mean(img,cycles,display=True,mean_only=False):
    img_set = img

    cs = cycles
    rows = 7
    time_array = np.zeros((rows,cycles),dtype = np.float32)
    mse_array = np.zeros((rows,cycles),dtype = np.float32)
    ssim_array = np.zeros((rows,cycles),dtype = np.float32)

    path = os.environ.get('TRAINING_DIR')+'/no_qat/checkpoint.pth'
    eval_image_set(img_set, path, False,'none', True)

    for i in range(cs):
        if (mean_only == False):
            print("==> experiment %2d" %i)
            t = PrettyTable(['training','quantization','time','mse','ssim'])
            t.float_format=".5"
            t.float_format['time'] = ".2"
            t.border=True
            t.align="r"
            t.align["quantization"]="l"
            t.align["training"]="l"

            path = os.environ.get('TRAINING_DIR')+'/no_qat/checkpoint.pth'
            mse_origin,ssim_origin = eval_origin(img_set, False)
            mse_noisy, ssim_noisy = eval_noisy(img_set, False)
            t.add_row(['NA (origin)','NA',"NA",mse_origin,ssim_origin])
            t.add_row(['NA (noisy)','NA',"NA",mse_noisy,ssim_noisy])

            eval_time_none,q, mse, ssim = eval_image_set(img_set, path, False,
e,'none', display)
            time_array[0,i] = (eval_time_none/eval_time_none)*100
            mse_array[0,i] = mse
            ssim_array[0,i] = ssim
            t.add_row(['standard',q,(eval_time_none/eval_time_none)*100,mse,
ssim])

            eval_time,q, mse, ssim = eval_image_set(img_set, path, False,'fx
_static', False)
            time_array[1,i] = (eval_time/eval_time_none)*100
            mse_array[1,i] = mse
            ssim_array[1,i] = ssim
            t.add_row(['standard',q,(eval_time/eval_time_none)*100,mse,ssim
])

            eval_time,q, mse, ssim = eval_image_set(img_set, path, False,'fx
_dynamic', False)
            time_array[2,i] = (eval_time/eval_time_none)*100
            mse_array[2,i] = mse
            ssim_array[2,i] = ssim
            t.add_row(['standard',q,(eval_time/eval_time_none)*100,mse,ssim
])

            path = os.environ.get('TRAINING_DIR')+'/qat/checkpoint.pth'
            eval_time,q, mse, ssim = eval_image_set(img_set, path, False,'no

```

```

ne', False)
    time_array[3,i] = (eval_time/eval_time_none)*100
    mse_array[3,i] = mse
    ssim_array[3,i] = ssim
    t.add_row(['QAT',q,(eval_time/eval_time_none)*100,mse,ssim])

    eval_time,q, mse, ssim = eval_image_set(img_set, path, True,'fx_
static', False)
    time_array[4,i] = (eval_time/eval_time_none)*100
    mse_array[4,i] = mse
    ssim_array[4,i] = ssim
    t.add_row(['QAT',q,(eval_time/eval_time_none)*100,mse,ssim])

    eval_time,q, mse, ssim = eval_image_set(img_set, path, True,'fx_
dynamic', False)
    time_array[5,i] = (eval_time/eval_time_none)*100
    mse_array[5,i] = mse
    ssim_array[5,i] = ssim
    t.add_row(['QAT',q,(eval_time/eval_time_none)*100,mse,ssim])

    path = os.environ.get('TRAINING_DIR')+'/qat_fx_static/checkpoint
t.pth'
#    eval_time,q, mse, ssim = eval_image_set(img_set, path, False,'non
e', False)
#    t.add_row(['QAT_fc_static',q,(eval_time/eval_time_none)*100,mse,ssi
m])
    eval_time,q, mse, ssim = eval_image_set(img_set, path, True,'fx_
static', False)
    time_array[6,i] = (eval_time/eval_time_none)*100
    mse_array[6,i] = mse
    ssim_array[6,i] = ssim
    t.add_row(['QAT_fc_static',q,(eval_time/eval_time_none)*100,mse,
ssim])
#    eval_time,q, mse, ssim = eval_image_set(img_set, path, True,'fx_dyn
amic', False)
#    t.add_row(['QAT_fc_static',q,(eval_time/eval_time_none)*100,mse,ssi
m])

    if (mean_only ==False):
        print(t)

    print("Summary of experiments")

    t = PrettyTable(['training','quantization','mean time','mean mse','m
ean ssim'])
    t.float_format['mean time'] = ".2"
    t.float_format['mean mse'] = ".5"
    t.float_format['mean ssim'] = ".5"
    t.border=True
    t.align="r"
    t.align["quantization"]="l"
    t.align["training"]="l"

    t.add_row(['standard','none',np.mean(time_array[0,:]),np.mean(mse_ar
ray[0,:]),np.mean(ssim_array[0,:])])
    t.add_row(['standard','fx_static',np.mean(time_array[1,:]),np.mean(m
se_array[1,:]),np.mean(ssim_array[1,:])])
    t.add_row(['standard','fx_dynami',np.mean(time_array[2,:]),np.mean(m


```



```
se_array[2,:]),np.mean(ssim_array[2,:]))
    t.add_row(['QAT', 'none', np.mean(time_array[3,:]), np.mean(mse_array[3
,:]), np.mean(ssim_array[3,:])])
    t.add_row(['QAT', 'fx_static', np.mean(time_array[4,:]), np.mean(mse_ar
ray[4,:]), np.mean(ssim_array[4,:])])
    t.add_row(['QAT', 'fx_dynamic', np.mean(time_array[5,:]), np.mean(mse_a
rray[5,:]), np.mean(ssim_array[5,:])])
    t.add_row(['QAT_fc_static', 'fx_static', np.mean(time_array[6,:]), np.m
ean(mse_array[6,:]), np.mean(ssim_array[6,:])])
    print(t)
```

```
In [11]: dataset_root_dir = os.environ.get('DATA_DIR')+'/images'
test_set = NoisyBSDSDataset(dataset_root_dir, mode='test', image_size=(3
20, 320))
img_set = test_set[1]
```

```
In [12]: experiment_mean(img_set,10,False, False)
```



==> experiment 0

```
/usr/local/lib/python3.6/dist-packages/torch/quantization/observer.py:1
24: UserWarning: Please use quant_min and quant_max to specify the rang
e for observers.                reduce_range will be deprecated in
a future release of PyTorch.
```

reduce\_range will be deprecated in a future release of PyTorch."

```
/usr/local/lib/python3.6/dist-packages/torch/_tensor.py:575: UserWarnin
g: floor_divide is deprecated, and will be removed in a future version
of pytorch. It currently rounds toward 0 (like the 'trunc' function NOT
'floor'). This results in incorrect rounding for negative values.
```

To keep the current behavior, use `torch.div(a, b, rounding_mode='trunc')`, or for actual floor division, use `torch.div(a, b, rounding_mode='floor')`. (Triggered internally at `/pytorch/aten/src/ATen/native/BinaryOps.cpp:467`.)

```
return torch.floor_divide(self, other)
```

```
/usr/local/lib/python3.6/dist-packages/torch/quantization/fx/quantizati
on_patterns.py:444: UserWarning: dtype combination: (torch.float32, tor
ch.qint8, torch.quint8) is not supported by Conv supported dtype combin
ations are: [(torch.quint8, torch.qint8, None)]
```

"supported dtype combinations are: {}".format(dtypes, supported\_dtypes))

training	quantization	time	mse	ssim
NA (origin)	NA	NA	0.00000	1.00000
NA (noisy)	NA	NA	0.01093	0.48157
standard	none	100.00	0.00158	0.87032
standard	fx_static	69.72	0.00160	0.86841
standard	fx_dynamic	105.88	0.00158	0.87032
QAT	none	92.25	0.00140	0.88876
QAT	fx_static	61.35	0.00142	0.88710
QAT	fx_dynamic	92.21	0.00140	0.88876
QAT_fc_static	fx_static	70.19	0.00428	0.68093

==> experiment 1

training	quantization	time	mse	ssim
NA (origin)	NA	NA	0.00000	1.00000
NA (noisy)	NA	NA	0.01093	0.48157
standard	none	100.00	0.00158	0.87032
standard	fx_static	66.26	0.00160	0.86841
standard	fx_dynamic	92.42	0.00158	0.87032
QAT	none	99.33	0.00140	0.88876
QAT	fx_static	69.78	0.00142	0.88710
QAT	fx_dynamic	97.02	0.00140	0.88876
QAT_fc_static	fx_static	65.67	0.00428	0.68093

==> experiment 2

training	quantization	time	mse	ssim
NA (origin)	NA	NA	0.00000	1.00000
NA (noisy)	NA	NA	0.01093	0.48157
standard	none	100.00	0.00158	0.87032
standard	fx_static	78.10	0.00160	0.86841
standard	fx_dynamic	111.15	0.00158	0.87032
QAT	none	111.84	0.00140	0.88876
QAT	fx_static	67.81	0.00142	0.88710
QAT	fx_dynamic	101.75	0.00140	0.88876
QAT_fc_static	fx_static	68.00	0.00428	0.68093

==> experiment 3

training	quantization	time	mse	ssim
NA (origin)	NA	NA	0.00000	1.00000
NA (noisy)	NA	NA	0.01093	0.48157
standard	none	100.00	0.00158	0.87032
standard	fx_static	64.06	0.00160	0.86841
standard	fx_dynamic	80.66	0.00158	0.87032
QAT	none	91.82	0.00140	0.88876
QAT	fx_static	71.79	0.00142	0.88710
QAT	fx_dynamic	84.86	0.00140	0.88876
QAT_fc_static	fx_static	68.12	0.00428	0.68093

==> experiment 4

training	quantization	time	mse	ssim
NA (origin)	NA	NA	0.00000	1.00000
NA (noisy)	NA	NA	0.01093	0.48157
standard	none	100.00	0.00158	0.87032
standard	fx_static	67.41	0.00160	0.86841
standard	fx_dynamic	91.06	0.00158	0.87032
QAT	none	95.97	0.00140	0.88876
QAT	fx_static	72.21	0.00142	0.88710
QAT	fx_dynamic	91.09	0.00140	0.88876
QAT_fc_static	fx_static	71.90	0.00428	0.68093

==> experiment 5

training	quantization	time	mse	ssim
NA (origin)	NA	NA	0.00000	1.00000
NA (noisy)	NA	NA	0.01093	0.48157
standard	none	100.00	0.00158	0.87032
standard	fx_static	68.97	0.00160	0.86841
standard	fx_dynamic	95.46	0.00158	0.87032
QAT	none	98.17	0.00140	0.88876
QAT	fx_static	64.60	0.00142	0.88710
QAT	fx_dynamic	86.53	0.00140	0.88876
QAT_fc_static	fx_static	68.79	0.00428	0.68093

==> experiment 6

training	quantization	time	mse	ssim
NA (origin)	NA	NA	0.00000	1.00000
NA (noisy)	NA	NA	0.01093	0.48157
standard	none	100.00	0.00158	0.87032
standard	fx_static	82.65	0.00160	0.86841
standard	fx_dynamic	95.64	0.00158	0.87032
QAT	none	103.51	0.00140	0.88876
QAT	fx_static	65.10	0.00142	0.88710
QAT	fx_dynamic	100.44	0.00140	0.88876
QAT_fc_static	fx_static	61.58	0.00428	0.68093

==> experiment 7

training	quantization	time	mse	ssim
NA (origin)	NA	NA	0.00000	1.00000
NA (noisy)	NA	NA	0.01093	0.48157
standard	none	100.00	0.00158	0.87032
standard	fx_static	61.83	0.00160	0.86841
standard	fx_dynamic	94.23	0.00158	0.87032
QAT	none	102.91	0.00140	0.88876
QAT	fx_static	62.22	0.00142	0.88710
QAT	fx_dynamic	93.65	0.00140	0.88876
QAT_fc_static	fx_static	62.16	0.00428	0.68093

==> experiment 8

training	quantization	time	mse	ssim
----------	--------------	------	-----	------

NA (origin)	NA	NA	0.00000	1.00000
NA (noisy)	NA	NA	0.01093	0.48157
standard	none	100.00	0.00158	0.87032
standard	fx_static	61.88	0.00160	0.86841
standard	fx_dynamic	101.99	0.00158	0.87032
QAT	none	106.04	0.00140	0.88876
QAT	fx_static	70.97	0.00142	0.88710
QAT	fx_dynamic	97.92	0.00140	0.88876
QAT_fc_static	fx_static	75.35	0.00428	0.68093

==> experiment 9

training	quantization	time	mse	ssim
NA (origin)	NA	NA	0.00000	1.00000
NA (noisy)	NA	NA	0.01093	0.48157
standard	none	100.00	0.00158	0.87032
standard	fx_static	65.97	0.00160	0.86841
standard	fx_dynamic	102.94	0.00158	0.87032
QAT	none	95.11	0.00140	0.88876
QAT	fx_static	57.19	0.00142	0.88710
QAT	fx_dynamic	78.97	0.00140	0.88876
QAT_fc_static	fx_static	61.48	0.00428	0.68093

Summary of experiments

training	quantization	mean time	mean mse	mean ssim
standard	none	100.0	0.0015840972	0.8703171
standard	fx_static	68.684296	0.0016042447	0.86840993
standard	fx_dynami	97.14357	0.0015840972	0.8703171
QAT	none	99.692245	0.0014023345	0.8887638
QAT	fx_static	66.30228	0.0014186573	0.8871044
QAT	fx_dynamic	92.44427	0.0014023344	0.8887638
QAT_fc_static	fx_static	67.32369	0.0042808643	0.6809325

The experiment clearly showed the significant processing time improvement without much of the quality degradation using all variants of `fx_static` quantization. Surprisingly, Quantization Aware Training did not produce superior results over `fx_static` method, and in some cases the quality were reduced. It is possible that QAT will depend on the type of the data and with different dataset or algorithm will produce significantly better results. That assumption needs to be explored. The field is still very new. The project used PyTorch latest quantization techniques and software, some of which is still in pre-release. PyTorch quantization support only x86 and ARM architectures. x86 backend was used in the project.

The testing was done with one original image and one noisy derivated image. The time fluctuation between experiments can be attributed to the compute infrastructure and cluster resources. The summary values appeared to be consistent over multiple experement runs performed.